

CS 33: Introduction to Computer Organization

Week 8

Admin

- Homework 5: “due” Wednesday, November 25th
- Lab 4: OpenMP (AKA your reward for finishing Smashing Lab) “due” Thursday, December 3rd.
- Lab 4: OpenMP (AKA your reward for finishing Smashing Lab) **due** Friday, December 4th.
- Yes, actually, honest to goodness **due** on Friday. That is the DROP DEAD DATE. Any assignments submitted beyond that date will NOT be accepted.

Agenda

- Midterm 2
- Wrap up Processes/Threads and Synchronization
- Lab 4: OpenMP
- HW5/Class Material included in the slides that there's no way in hell we'll get to:
 - Command line args
 - Deadlock Cases
 - Thread safe vs. Reentrancy

Midterm 2

- Mean: 48.1 (out of 99)
- Median: 48 (out of 99)
- Interestingly in my section, something of a bimodal distribution.

Midterm Secret Solutions

- Question 1: Though unconventional, I insist on me

REDACTED

eyeball of
the unusual number of gorillas.

Processes

- What happens when a child process completes?
- Processes can be “running”, “stopped”, or “terminated”.
- A completed child process is terminated, but the resources that are used to keep track of the child process still exist.
- This is (and this is completely real), known as a zombie. A zombie child.

Processes

- An undead process is one who has technically terminated.
 - More formal definition of `exit(0)`: terminates the process
- However, you can't get rid of a zombie that easily; the resources that the operating system uses to manage that process continue to exist.
- A process will do so unless, and I'm dead serious, it is “reaped”, after which its resources are reclaimed.

Processes

- One of the functions that the parent has over the child is:
 - `waitpid(pid_t pid, int *status, int options)`
- From the parent's perspective, the default behavior is “wait for the child process with process id of 'pid' to terminate before continuing”
- If `pid == -1`, `waitpid` will wait for any one of the child processes to complete.

Processes

- However, waitpid technically serves an additional role regarding system resources.
- waitpid essentially waits for the child to die and become a zombie so that the parent can reap it of its resources.
- ...which is a terrifying glimpse into the disturbed mind of whoever came up with this.

Processes

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
int main()
{
    if(fork() == 0)
    {
        printf("a");
    }
    else
    {
        printf("b");
        waitpid(-1, NULL, 0);
    }
    printf("c");
    exit(0);
}
```

- What's the output of this program?

Processes

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
int main()
{
    if(fork() == 0)
    {
        printf("a");
    }
    else
    {
        printf("b");
        waitpid(-1, NULL, 0);
    }
    printf("c");
    exit(0);
}
```

- Once fork is hit, the child will print a and the parent will print b.
- Afterwards, both will print c.
- The parent will not print c until the child has printed a and c.
- There are no other restrictions on ordering.
- abcc, acbc, bacc

Processes

- Because variables that are created before a fork is duplicated in the child process, this means that child processes also share file descriptors.
- The OS needs to be able to keep track of what files that each process has opened.
- The OS does this by maintaining a table of processes and their corresponding files that are open.
- Each entry has a label or a “file descriptor”

Processes


- In order for the process to access the file, it has to go through the OS so it must know the file descriptor of the file it is trying to access.
- A file descriptor is just an integer that corresponds to an index into the OS's table of files.
- Thus, children also have duplicates of the file descriptors and can thus access the same files.

Processes

- File descriptors have to be closed in order to tell the OS that it's done dealing with a particular file.
- As a result, when a child gets a cloned file descriptor, the parent and child are responsible for closing the file descriptor.

Processes

- With Processes, the organization looks like this:
- If a child process becomes orphaned (ex. process 1 exits while process 4 and process 5 are still running), the OS arranges for the process to be adopted by the kindly Mrs. Init Process from down the street.



Processes

- Because of this tree hierarchy and because most processes run in isolation (no shared memory), Processes really only have to worry about the high level status of the children:
 - Is it still running? Paused?
 - How did it exit (normally or via signal)?
 - Why was it out so late last night? Has it been drinking? Was it Braden? You know he's a bad kid.
- As a result, processes don't have to worry much about synchronization. Not so with the hyper-social Threads.

Threads

- All threads that belong to a particular process are peers, which mean there's a lot more management that can occur.
- `int pthread_create(pthread_t *tid, pthread_attr_t const *attr, func *f, void *args)`
 - Creates thread, assigns thread id to tid. When thread starts, it will call `f(args)`
- `pthread_t pthread_self(void);`
 - Get own thread id.
- `void pthread_exit(void *thread_return)`
 - Thread kills itself with return value in `*thread_return`

Threads

- `int pthread_cancel(pthread_t tid)`
 - Thread murders another thread with id `tid`, in cold blood.
- `int pthread_join(pthread_t tid, void **thread_return)`
 - One thread waits for thread `tid` to complete before continuing.
 - Technically more reaping happens here.

Threads

- `int pthread_detach(pthread_t tid)`
 - Makes thread specified by `tid` autonomous. That is, the thread cannot be killed via `pthread_cancel` and it cannot be reaped/waited on by `pthread_join`.
 - The thread achieves invincibility.

Threads

Consider the following code:

```
int main()
{
    while(true)
    {
        pthread_t tid;
        int connfd = Accept(...);
        pthread_create(&tid, 0,
thread, &connfd);
    }
}
```

```
void *thread(void * args)
{
    pthread_detach(pthread_self());
    int * pconnfd = args;
    int connfd = *pconnfd;
    process(connfd);
    close(connfd);
    return 0;
}
```

- Does this work?

Threads

- connfd is a local variable of the main thread.
- Thus, connfd has an address that corresponds to some spot on the main.
- The value of connfd is assigned in the main thread by Accept.
- The address to this value of connfd is passed into the thread.

```
while(true)
{
    pthread_t tid;
    int connfd = Accept(...);
    pthread_create(&tid, 0,
thread, &connfd);
}
```

```
void *thread(void * args)
{
    ...
    int * pconnfd = args;
    int connfd = *pconnfd
    ...
}
```

Threads

- The thread then uses the address of the main thread's `connfd` for args.
- The thread will get the value and store it to a local variable.
- Meanwhile, the main thread re-assigns a value to `connfd`.
- The address that the old thread uses still corresponds to `connfd` in the main thread.

```
while(true)
{
    pthread_t tid;
    int connfd = Accept(...);
    pthread_create(&tid, 0,
thread, &connfd);
}
```

```
void *thread(void * args)
{
    ...
    int * pconnfd = args;
    int connfd = *pconnfd
    ...
}
```

Threads

- If that reassignment (connfd = Accept...) occurs before connfd = *pconnfd, then the old value will be wiped away.
- This is an example of a race condition (and a pretty bad one).
- Make sure this doesn't happen.

```
while(true)
{
    pthread_t tid;
    int connfd = Accept(...);
    pthread_create(&tid, 0,
thread, &connfd);
}
```

```
void *thread(void * args)
{
    ...
    int * pconnfd = args;
    int connfd = *pconnfd
    ...
}
```

Threads

- Conclusion?
- Any time a thread gets a pointer that points to another thread's stack frame, be wary.
- We can fix this by making it so that the value passed into the thread is NOT a pointer pointing to the main thread's stack frame.

```
while(true)
{
    pthread_t tid;
    int connfd = Accept(...);
    pthread_create(&tid, 0,
thread, &connfd);
}
```

```
void *thread(void * args)
{
    ...
    int * pconnfd = args;
    int connfd = *pconnfd
    ...
}
```


Threads

- Now, we deal with an `int *`, whose address IS still in the main thread's stack.
- However, what we pass into the thread is the value of the `int*`, which is malloc'd.

```
while(true)
{
    pthread_t tid;
    int *connfd = malloc(sizeof(int));
    *connfd = Accept(...);
    pthread_create(&tid, 0, thread,
connfd);
}
void *thread(void * args)
{
    ...
    int * pconnfd = args;
    int connfd = *pconnfd
    ...
}
```

Synchronization + Race Conditions

- One of the major advantages of using threads over processes is the ease by which threads can share data.
- This of course, is marred by the potential that by using threads, you get the wrong result.
- That's bad.
- How do we ensure that behavior is correct, even when the execution order is not guaranteed?

Synchronization + Race Conditions

- We can enforce protection for shared variables or critical sections with semaphores, which are basically locks that tell us how many threads can enter a section of code.
- The semaphore is of type: `sem_t` and it is sort of a glorified counter or more conceptually, a door that allows/prevents entry into a section of code.

Synchronization + Race Conditions

- While the counter is non-zero, the door is open and thread can pass. When this happens, the counter decrements.
- When the counter is zero, the door is closed. The thread must wait for the door to be open again (counter becomes non-zero) before it can pass.

Synchronization + Race Conditions

- `sem_t sem;`
- `sem_wait(&sem)` – If `sem` is non-zero, return and decrement `sem` (the door is open, thread is allowed to pass). If `sem` is zero, wait until `sem` is non-zero, then decrement and return (the door is closed, wait for it to open).
- `sem_post(&sem)` – Increment `sem` by one. If `sem` was previously zero, the door was closed. Open the door and allow another thread in.

Synchronization + Race Conditions

- Say we have the following line of code:
 - <line of code>
- ...and we want to allow only two threads to be able to execute that line of code at any time.
- `sem_t sem;`
- `sem_init(&sem, 0, 2);` //0 is an options argument, 2 is the number of threads allowed
- `sem_wait(&sem);`
- <line of code>;
- `sem_post(&sem);`

Synchronization + Race Conditions

- Let's say we have some shared resource that you can read and write from. This can be done via two functions:

```
void read()  
{  
    <read shared resource>  
}
```

```
void write()  
{  
    <write shared resource>  
}
```

Synchronization + Race Conditions

- However, we want the following conditions:
 - There can be an unbounded number of concurrent readers.
 - There can be only 1 writer and if someone is writing, there can be no readers.
- How can we go about this?

Synchronization + Race Conditions

```
void read()  
{  
    <read>  
}
```

```
void write()  
{  
    <write>  
}
```

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
void read()
{
    sem_wait(&r);
    <read>
    sem_post(&r);
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- We'll definitely want some semaphores around the critical sections in read/write.
- Is this right?

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
void read()
{
    sem_wait(&r);
    <read>
    sem_post(&r);
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- This prevents multiple writers, but it allows writing to happen at the same time as reading.
- Also, what do we initialize r to
We want unbounded readers.
- Let's try to solve the
“writing/reading at the same time” problem first

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
void read()
{
    sem_wait(&r)
    sem_wait(&w)
    <read>
    sem_wait(&w)
    sem_wait(&r)
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- Now, writing cannot happen at the same time as reading.
- What's the next step?

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
void read()
{
    sem_wait(&r)
    sem_wait(&w)
    <read>
    sem_wait(&w)
    sem_wait(&r)
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- If we initialized r to, say 10, we could not have 10 readers reading at the same time because w is initialized to 1.
- We recognize that we only want to run `sem_wait(&w)` when there are no readers.

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
void read()
{
    sem_wait(&r)
    sem_wait(&w)
    <read>
    sem_wait(&w)
    sem_wait(&r)
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- If no readers, there might be a writer. Therefore, a reader must call `sem_wait(&w)` to reserve the critical section.
- If there are readers, then that means there cannot be a writer.
- Therefore, the reader doesn't need to call `sem_wait(&w)`

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
void read()
{
    sem_wait(&r)
    if(there are no readers)
        sem_wait(&w)

    <read>

    if(this is the last reader)
        sem_post(&w)
    sem_post(&r)
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- Once a reader has finished reading, it must also check: if it is the last reader, then it should be able to allow a writer. Therefore it must release the lock by calling `sem_post(&w)`

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
int reader_count = 0;
void read()
{
    sem_wait(&r)
    reader_count++;
    if(reader_count == 1)
        sem_wait(&w)
```

<read>

```
    reader_count--;
    if(reader_count == 0)
        sem_post(&w)
    sem_post(&r)
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- Does this work?

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
int reader_count = 0;
void read()
{
    sem_wait(&r)
    reader_count++;
    if(reader_count == 1)
        sem_wait(&w)

    <read>

    reader_count--;
    if(reader_count == 0)
        sem_post(&w)
        sem_post(&r)
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- Two problems, first of all, this only works if we can set r to infinity, which we can't.
- Second, there's still a possible race condition. What if two reader threads increment reader_count before either can get to reader_count == 1? The whole system messes up.

Synchronization + Race Conditions

```
sem_t r; // init to ?
sem_t w; // init to 1
int reader_count = 0;
void read()
{
    sem_wait(&r)
    reader_count++;
    if(reader_count == 1)
        sem_wait(&w)

    <read>

    reader_count--;
    if(reader_count == 0)
        sem_post(&w)
    sem_post(&r)
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- That this treats the <read> as the critical section, but that's not what we want to protect.
- We want to make sure that infinite readers can read, but only one reader can increment reader_count at a time.

Synchronization + Race Conditions

```
sem_t r; // init to 1
sem_t w; // init to 1
int reader_count = 0;
void read()
{
    sem_wait(&r);
    reader_count++;
    if(reader_count == 1)
        sem_wait(&w);
    sem_post(&r);
    <read>
    sem_wait(&r);
    reader_count--;
    if(reader_count == 0)
        sem_post(&w);
    sem_post(&r);
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- It satisfies the constraints (single writer, multiple reader).
- Is there maybe still a problem though?

Synchronization + Race Conditions

```
sem_t r; // init to 1
sem_t w; // init to 1
int reader_count = 0;
void read()
{
    sem_wait(&r);
    reader_count++;
    if(reader_count == 1)
        sem_wait(&w);
    sem_post(&r);
    <read>
    sem_wait(&r);
    reader_count--;
    if(reader_count == 0)
        sem_post(&w);
    sem_post(&r);
}
```

```
void write()
{
    sem_wait(&w);
    <write>
    sem_post(&w);
}
```

- This method is unfair to writers.
- That is to say, a potential writer could be starved since if readers keep coming in, the writer will never have a chance to write.

Lab 4: OpenMP Lab

- Unlike the other labs, the objective of this lab is simple.
- Premise: You have code that's slow.
- Goal: Make it go fast.
- The End.

Lab 4: OpenMP Lab

- Getting started:
 - Download the openmplab.tgz file from CCLE
 - Copy it to Inxsrv09.seas.ucla.edu (more than ever, USE Inxsrv09 and gcc version 5.2.0)
 - Unzip it.
- To run the unaltered code, use:
 - make seq
 - This will produce an executable called “seq” (note, the specs are wrong).

Lab 4: OpenMP Lab

- In order to speed it up, you are to use two tools, one familiar and one new (probably):
 - To determine what to speed up: gprof.
 - To speed it up: OpenMP, an API for shared memory, thread based parallelism.

Lab 4: OpenMP Lab

- To use gprof:
- Compile with GPROF=1
 - ex. make seq GPROF=1
- Execute the executable (this produces gmon.out)
 - ex. ./seq
- Run gprof with the executable as the argument
 - ex. gprof seq | less
 - The command gprof seq normally prints out a wall of text. The “| less” redirects the output as input into the “less” command, which allows you to read it more easily.
 - Quit “less” with q.

Lab 4: OpenMP Lab

- The code provided performs a computation using six functions (func0 to func5) located in the func.c.
- You are to speed up the execution of the program by speeding up these six function.
- How?

Lab 4: OpenMP Lab

- You could use the normal Chapter 5 style optimizations, but we want something powerful that makes uses of multiple processors.
- This means using threads and doing multithreading.
- However, we want to use something a little more powerful and exotic than POSIX threads.
- So exotic that you might never see it again...

Lab 4: OpenMP Lab

- As a result, we use OpenMP
- OpenMP is an extension to C/C++ (and Fortran) that allows for a simple and convenient way of launching threads to do work in parallel among multiple processors.
- Surprisingly, you won't be personally using a bunch of new functions.
- You'll be specifying a bunch of pre-processor directives.

Lab 4: OpenMP Lab

- Yup, these: “#”.
- As a refresher, preprocessor directives are statements that are marked with “#”.
- Statements marked with these are instructions that tell the processor what to do:
 - `#define min(a, b) (a < b ? a : b)`
 - `#include <stdio.h>`
 - `#don't become sentient`

Lab 4: OpenMP Lab

- The first step of compilation is the preprocessing step where the compiler will find all of the directives and follow the instructions.
- OpenMP exposes an interface for multithreading that is almost purely based on these preprocessor instructions.
- Other than some helper functions defined in `omp.h`. Ex.
 - `omp_set_num_threads(int);`
 - `omp_get_thread_num();`

Lab 4: OpenMP Lab

- Ex:

```
int main()
{
    int a = 0;
    printf("%d\n", a);
}
```

- Say you want to parallelize this so that 4 threads run this code.

Lab 4: OpenMP Lab

Ex:

```
int main()
{
    omp_set_num_threads(4); //Use 4 threads
    #pragma omp parallel //Define the following
    {                               //block as parallel code
        int a = 0; //run by each thread
        printf("%d\n", a);
    }
}
```

Lab 4: OpenMP Lab

Ex:

```
int main()
{
    #pragma omp parallel
    {
        int a = 0;
        printf("%d\n", a);
    }
}
```

- If the number of threads is not specified, the default is used. On SEASnet Inxsrv09, this means 32 (the number of processors).

Lab 4: OpenMP Lab

- Upon hitting the block specified by the `#pragma omp` directive, the single thread of execution branches off into several parallel threads.
- When each thread reaches the end of the block:

```
#pragma omp parallel  
{  
    ...  
} ← Here
```
- ...each thread will wait for all of the other threads before continuing. This is known as a “barrier”.

Lab 4: OpenMP Lab

- The “parallel” keyword means “launch threads in parallel”
- This the basis how OpenMP works.
- Synchronization and control can also be done via these directives. Within a parallel block, you can also have other directives that specify things.
- Such as...

Lab 4: OpenMP Lab

- critical:

```
#pragma omp critical
{
    <CRITICAL SECTION>
}
```

- Only one thread can enter the critical section at a time. Other threads must wait (like semaphores)

- atomic:

```
#pragma omp atomic
{
    tmp += 1;
}
```

- Only one thread can perform a particular read/write. This only applies to single variable that is read/write.

Lab 4: OpenMP Lab

- barrier:
 - `#pragma omp barrier`
 - All threads must wait here for all threads to complete before continuing.
- ...and so much more.
- So... so much more.
- <http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>

Lab 4: OpenMP Lab

- Back to this example:

```
int main()
{
    #pragma omp parallel
    {
        int a = 0;
        printf("%d\n", a);
    }
}
```

- What if a were declared outside of the parallel section?

Lab 4: OpenMP Lab

- Back to this example:

```
int main()
{
    int a = 0;
    #pragma omp parallel
    {
        a++;
        printf("%d\n", a);
    }
}
```

- Is 'a' shared among the 32 threads?

Lab 4: OpenMP Lab

- Data sharing:
 - By default, a variable declared inside of a parallel block is private to each block.
 - By default, a variable declared outside of a parallel block is shared among all of the threads of a parallel block that accesses the variable. In the previous example, “a” was the same “a” among all threads.
 - “a” was also the same “a” that was declared before entering the parallel block
 - What do we expect to see as an output?

Lab 4: OpenMP Lab

- Data sharing:
 - Utter chaos
 - Because there is no protection of the shared variable and because writes are not atomic, the results are going to be crazy.
 - If we want to share the variables in a way that's more graceful and predictable, we'll have to use critical or atomic but then, we'll essentially be doing code in serial.

Lab 4: OpenMP Lab

- A variable within a parallel block can be made private to each thread by using the “private” keyword. Note: this example is deliberately problematic. To see why, turn the page...

```
int main()
{
    int a = 0;
    #pragma omp parallel private(a)
    {
        a++;
        printf("%d\n", a);
    }
}
```

Lab 4: OpenMP Lab

- Warning: the variable will also be private relative to the original declaration.

```
int main()
{
    int a = 123;
    #pragma omp parallel private(a)
    {
        a++;
        printf("%d\n", a);
    }
}
```

- The “a” inside the block will not be initialized to 123. It will be uninitialized.

Lab 4: OpenMP Lab

- If you do want to make the private variable in the block initialized to the value specified in the original thread, use “firstprivate”.

```
int main()
{
    int a = 123;
    #pragma omp parallel firstprivate(a)
    {
        a++;
        printf(“%d\n”, a);
    }
}
```

- The “a” inside the block will now be initialized to 123.

Lab 4: OpenMP Lab

- Loops
- Consider:

```
int i;  
for(i = 0; i < N; i++)  
{  
    <PARALLELIZABLE CODE>  
}
```

- You can do this with an “#pragma omp parallel” block, but you'll have to manually assign threads to iterations of the loop.

Lab 4: OpenMP Lab

- Loops
- Or you can do this:

```
int i;
```

```
#pragma omp parallel for
```

```
for(i = 0; i < N; i++)
```

```
{
```

```
    <PARALLELIZABLE CODE>
```

```
}
```

- ...which magics it up for you.
- However, there are (at least) two bits of trickiness here.

Lab 4: OpenMP Lab

- What if you have this:

```
int acc = 0;
int i;
#pragma omp parallel for
for(i = 0; i < N; i++)
{
    acc += <PARALLELIZABLE CODE>
}
printf("%d\n", acc);
```

- What's the problem with this?

Lab 4: OpenMP Lab

- What if you have this:

```
int acc = 0;
int i;
#pragma omp parallel for
for(i = 0; i < N; i++)
{
    acc += <PARALLELIZABLE CODE>
}
printf("%d\n", acc);
```

- What's the problem with this?
 - `acc += ...` is not atomic. The result will not be correct.

Lab 4: OpenMP Lab

- To ensure that you can use accumulators use the “reduction” keyword, which allows you to collect results into one shared variable among threads.

```
int acc = 0;
int i;
#pragma omp parallel for reduction(+:acc)
for(i = 0; i < N; i++)
{
    acc += <PARALLELIZABLE CODE>
}
printf(“%d\n”, acc);
```

- Format: (<operation> : <variable>);

Lab 4: OpenMP Lab

- Finally, consider:

```
int acc = 0;
int j = 0;
int i;
#pragma omp parallel for reduction(+:acc)
for(i = 0; i < N; i++)
{
    j += 2;
    acc += <PARALLELIZABLE CODE>
}
printf("%d\n", acc);
```

- It looks like j is in a similar situation as acc.

Lab 4: OpenMP Lab

- However, the dependence on j is artificially loop carried. You can convert it to:

```
int acc = 0;
int j = 0;
int i;
#pragma omp parallel for private(j) reduction(+:acc)
for(i = 0; i < N; i++)
{
    j = 2*i;
    acc += <PARALLELIZABLE CODE>
}
printf("%d\n", acc);
```

- You may find that small tweaks like this may be necessary to make an “omp parallel for” possible.

Lab 4: OpenMP Lab

- Final Notes:
- The specs specify methods of invoking the make file, several of which are wrong.
- “make seq” will generate an executable called “seq”.
- “make omp” will generate an executable called “omp”.
- “make ____ SRC=funcv2.c” does not work.
 - If you'd like this functionality to work, in the Makefile change:
 - SRC = filter.c
 - SRCS = \$(SRC) main.c func.c util.c
 - to
 - SRC = func.c
 - SRCS = \$(SRC) main.c filter.c util.c

Lab 4: OpenMP Lab

- Final Notes:
 - The run times will vary each time you call filter.
 - As it turns out, as the SEASnet processor become more burdened, it will actually slow itself down, which will lead to a slower execution time.
 - If you think you have a solution, double check your time by running your code repeatedly.
 - You are supposed to speed up the “FUNC TIME” as reported by the print statement by 3.5x.

Lab 4: OpenMP Lab

- Final Notes:
 - When you run the code with GPROF=1, the times will slow down. Example output:
 FUNC TIME: .670985
 TOTAL TIME: 2.772028
 - This suggests that the total amount of time spent by the functions in func.c take up ~20-25%.
 - If you actually look at the gprof print out, you may see this:

Lab 4: OpenMP Lab

- Final Notes:

% time	cumulative seconds	self seconds	self calls	total ms/call	ms/call	name
61.19	0.44	0.44	15	29.37	31.65	func1
25.03	0.62	0.18	5177344	0.00	0.00	rand2
2.78	0.64	0.02	491520	0.00	0.00	f...
2.78	0.66	0.02	2	10.01	10.01	init
...						

- This suggests that func1 alone is taking 61% of the time.
- I'm no expert, but I think 20 is like, not the same as 61.

Lab 4: OpenMP Lab

- As it turns out, gprof doesn't take into account blocking I/O time, which is the bulk of what the non-func.c functions do.
- Use either your own measurements (get_time) or use the ones that are printed out by the output.

Lab 4: OpenMP Lab

- Final Notes: Memory Leaks
- According to checkmem, OpenMP loves to leak memory.

```
int a;  
#pragma omp parallel  
{  
    a = 0;  
}
```

- Even this “leaks memory”.

Lab 4: OpenMP Lab

- Final Notes:
- If the memory error printed out by “make checkmem” is “Memory not freed”, don't worry about it. Otherwise, go back and double check.
- Make sure to use “make check” to confirm that the output.txt generated by the program is the same as correct.txt. If silent, the files are the same.