



# A “Hands-on” Introduction to OpenMP\*

**Tim Mattson**


**Intel Corp.**

[timothy.g.mattson@intel.com](mailto:timothy.g.mattson@intel.com)

\* The name “OpenMP” is the property of the OpenMP Architecture Review Board.



# Outline


- **Unit 1: Getting started with OpenMP**
  - ♦ Mod1: Introduction to parallel programming
  - ♦ Mod 2: The boring bits: Using an OpenMP compiler (hello world)
  - ♦ Disc 1: Hello world and how threads work
- **Unit 2: The core features of OpenMP**
  - ♦ Mod 3: Creating Threads (the Pi program)
  - ♦ Disc 2: The simple Pi program and why it sucks
  -  ♦ Mod 4: Synchronization (Pi program revisited)
  - ♦ Disc 3: Synchronization overhead and eliminating false sharing
  - ♦ Mod 5: Parallel Loops (making the Pi program simple)
  - ♦ Disc 4: Pi program wrap-up
- **Unit 3: Working with OpenMP**
  - ♦ Mod 6: Synchronize single masters and stuff
  - ♦ Mod 7: Data environment
  - ♦ Disc 5: Debugging OpenMP programs
  - ♦ Mod 8: Skills practice ... linked lists and OpenMP
  - ♦ Disc 6: Different ways to traverse linked lists
- **Unit 4: a few advanced OpenMP topics**
  - ♦ Mod 8: Tasks (linked lists the easy way)
  - ♦ Disc 7: Understanding Tasks
  - ♦ Mod 8: The scary stuff ... Memory model, atomics, and flush (pairwise synch).
  - ♦ Disc 8: The pitfalls of pairwise synchronization
  - ♦ Mod 9: Threadprivate Data and how to support libraries (Pi again)
  - ♦ Disc 9: Random number generators
- **Unit 5: Recapitulation**

# OpenMP Overview:

## How do threads interact?


Recall our high level overview of OpenMP?

- OpenMP is a multi-threading, shared address model.
  - Threads communicate by sharing variables.
- Unintended sharing of data causes race conditions:
  - race condition: when the program's outcome changes as the threads are scheduled differently.
- To control race conditions:
  - Use synchronization to protect data conflicts.
- Synchronization is expensive so:
  - Change how data is accessed to minimize the need for synchronization.




# Synchronization:

- Synchronization: bringing one or more threads to a well defined and known point in their execution.
- The two most common forms of synchronization are:



**Barrier:** each thread wait at the barrier until all threads arrive.



**Mutual exclusion:** Define a block of code that only one thread at a time can execute.



# Synchronization

- **High level synchronization:**

- critical

- atomic

- barrier

- ordered

- **Low level synchronization**

- flush

- locks (both simple and nested)

**Synchronization is used  
to impose order  
constraints and to  
protect access to shared  
data**

Discussed  
Later



# Synchronization: Barrier

- **Barrier**: Each thread waits until all threads arrive.

```
#pragma omp parallel
{
    int id=omp_get_thread_num();
    A[id] = big_calc1(id);
    #pragma omp barrier


    B[id] = big_calc2(id, A);
}
```



# Synchronization: critical

- Mutual exclusion: Only one thread at a time can enter a **critical** region.

Threads wait  
their turn –  
only one at a  
time calls  
consume()



```
float res;  
  
#pragma omp parallel  
{   float B;  int i, id, nthrds;  
    id = omp_get_thread_num();  
    nthrds = omp_get_num_threads();  
    for(i=id;i<niters;i+=nthrds){  
        B = big_job(i);  
        #pragma omp critical  
        res += consume (B);  
    }  
}
```



# Synchronization: Atomic (basic form)

- **Atomic** provides mutual exclusion but only applies to the update of a memory location (the update of X in the following example)

```
#pragma omp parallel
```

```
{
```

```
    double tmp, B;
```

```
    B = DOIT();
```

```
    tmp = big_ugly(B);
```

```
#pragma omp atomic
```

```
    X += tmp;
```

```
}
```

The statement inside the atomic must be one of the following forms:

- $x \text{ binop} = \text{expr}$
- $x++$
- $++x$
- $x--$
- $--x$

X is an lvalue of scalar type and binop is a non-overloaded built in operator.

Additional forms of atomic were added in OpenMP 3.1.  
We will discuss these later.



## Exercise 3

- In exercise 2, you probably used an array to create space for each thread to store its partial sum.
- If array elements happen to share a cache line, this leads to false sharing.
  - Non-shared data in the same cache line so each update invalidates the cache line ... in essence “sloshing independent data” back and forth between threads.
- Modify your “pi program” from exercise 2 to avoid false sharing due to the sum array.

# Outline

- **Unit 1: Getting started with OpenMP**
  - ♦ Mod1: Introduction to parallel programming
  - ♦ Mod 2: The boring bits: Using an OpenMP compiler (hello world)
  - ♦ Disc 1: Hello world and how threads work
- **Unit 2: The core features of OpenMP**
  - ♦ Mod 3: Creating Threads (the Pi program)
  - ♦ Disc 2: The simple Pi program and why it sucks
  - ♦ Mod 4: Synchronization (Pi program revisited)
  - ♦ Disc 3: Synchronization overhead and eliminating false sharing
  - ♦ Mod 5: Parallel Loops (making the Pi program simple)
  - ♦ Disc 4: Pi program wrap-up
- **Unit 3: Working with OpenMP**
  - ♦ Mod 6: Synchronize single masters and stuff
  - ♦ Mod 7: Data environment
  - ♦ Disc 5: Debugging OpenMP programs
  - ♦ Mod 8: Skills practice ... linked lists and OpenMP
  - ♦ Disc 6: Different ways to traverse linked lists
- **Unit 4: a few advanced OpenMP topics**
  - ♦ Mod 8: Tasks (linked lists the easy way)
  - ♦ Disc 7: Understanding Tasks
  - ♦ Mod 8: The scary stuff ... Memory model, atomics, and flush (pairwise synch).
  - ♦ Disc 8: The pitfalls of pairwise synchronization
  - ♦ Mod 9: Threadprivate Data and how to support libraries (Pi again)
  - ♦ Disc 9: Random number generators
- **Unit 5: Recapitulation**

# Pi program with false sharing\*

- Original Serial pi program with 100000000 steps ran in 1.83 seconds.

## Example: A simple Parallel pi program

```
#include <omp.h>
static long num_steps = 100000;    double step;
#define NUM_THREADS 2
void main ()
{
    int i, nthreads; double pi, sum[NUM_THREADS];
    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int i, id, nthrds;
        double x;
        id = omp_get_thread_num();
        nthrds = omp_get_num_threads();
        if (id == 0) nthreads = nthrds;
        for (i=id, sum[id]=0.0; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            sum[id] += 4.0/(1.0+x*x);
        }
        for(i=0, pi=0.0; i<nthreads; i++) pi += sum[i] * step;
    }
}
```

**Recall that promoting sum to an array made the coding easy, but led to false sharing and poor performance.**

threads	1 <sup>st</sup> SPMD
1	1.86
2	1.03
3	1.08
4	0.97

\*Intel compiler (icpc) with no optimization on Apple OS X 10.7.3 with a dual core (four HW thread) Intel® Core™ i5 processor at 1.7 Ghz and 4 Gbyte DDR3 memory at 1.333 Ghz.

## Example: Using a critical section to remove impact of false sharing

```
#include <omp.h>
```

```
static long num_steps = 100000;    double step;
```

```
#define NUM_THREADS 2
```

```
void main ()
```

```
{    double pi;    step = 1.0/(double) num_steps;
```

```
    omp_set_num_threads(NUM_THREADS);
```

```
#pragma omp parallel
```

```
{
```

```
    int i, id, nthrds;    double x, sum,
```

```
    id = omp_get_thread_num();
```

```
    nthrds = omp_get_num_threads();
```

```
    if (id == 0) nthrds = nthrds;
```

```
    for (i=id, sum=0.0; i< num_steps; i=i+nthreads)
```

```
    {
```

```
        x = (i+0.5)*step;
```

```
        sum += 4.0/(1.0+x*x);
```

```
    }
```

```
#pragma omp critical
```

```
    pi += sum * step;
```

```
}  
}
```

Create a scalar local to each thread to accumulate partial sums.

No array, so no false sharing.

Sum goes "out of scope" beyond the parallel region ... so you must sum it in here. Must protect summation into pi in a critical region so updates don't conflict

# Results\*: pi program critical section

- Original Serial pi program with 100000000 steps ran in 1.83 seconds.

## Example: Using a critical section to remove impact of false sharing

```
#include <omp.h>
static long num_steps = 100000;    double step;
#define NUM_THREADS 2
void main ()
{
    double pi;    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int i, id, nthrds;    double x, sum;
        id = omp_get_thread_num();
        nthrds = omp_get_num_threads();
        if (id == 0) nthrds = nthrds;
        id = omp_get_thread_num();
        nthrds = omp_get_num_threads();
        for (i=id, sum=0.0; i< num_steps; i=i+nthrds){
            x = (i+0.5)*step;
            sum += 4.0/(1.0+x*x);
        }
        #pragma omp critical
        pi += sum * step;
    }
}
```

threads	1st SPMD	1st SPMD padded	SPMD critical
1	1.86	1.86	1.87
2	1.03	1.01	1.00
3	1.08	0.69	0.68
4	0.97	0.53	0.53

\*Intel compiler (icpc) with no optimization on Apple OS X 10.7.3 with a dual core (four HW thread) Intel® Core™ i5 processor at 1.7 Ghz and 4 Gbyte DDR3 memory at 1.333 Ghz.

## Example: Using a critical section to remove impact of false sharing

```
#include <omp.h>
static long num_steps = 100000;      double step;
#define NUM_THREADS 2
void main ()
{
    double pi;      step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
    int i, id, nthrds;
    double x;
    id = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    if (id == 0) nthrds = nthrds;
    for (i=id, sum=0.0; i< num_steps; i=i+nthreads)
    {
        x = (i+0.5)*step;
        #pragma omp critical
        pi += 4.0/(1.0+x*x);
    }
}
pi *= step;
}
```

**Be careful  
where you put  
a critical  
section**

What would happen if  
you put the critical  
section inside the loop?

## Example: Using an atomic to remove impact of false sharing

```
#include <omp.h>
```

```
static long num_steps = 100000;    double step;
```

```
#define NUM_THREADS 2
```

```
void main ()
```

```
{    double pi;    step = 1.0/(double) num_steps;
```

```
    omp_set_num_threads(NUM_THREADS);
```

```
#pragma omp parallel
```

```
{
```

```
    int i, id, nthrds;    double x, sum;
```

```
    id = omp_get_thread_num();
```

```
    nthrds = omp_get_num_threads();
```

```
    if (id == 0) nthrds = nthrds;
```

```
    for (i=id, sum=0.0; i< num_steps;
```

```
        i=i+nthrds)
```

```
    {
```

```
        x = (i+0.5)*step;
```

```
        sum += 4.0/(1.0+x*x);
```

```
    }
```

```
    sum = sum*step;
```

```
#pragma atomic
```

```
    pi += sum ;
```

```
}
```

Create a scalar local to each thread to accumulate partial sums.

No array, so no false sharing.

Sum goes “out of scope” beyond the parallel region ... so you must sum it in here. Must protect summation into pi so updates don't conflict